

Training the dataset

- Two type of dataset was created , dataset with stopword removal and lemmatization and dataset without stopword removal and lemmatization
- Dataset was trained on Bert and Roberta with 6 epoch

Feature	BERT-base	RoBERTa-base
Architecture	12 Transformer encoder layers, 768 hidden size, 12 self-attention heads	12 Transformer encoder layers, 768 hidden size, 12 self-attention heads
Parameters	~110 million	~125 million
Tokenizer	WordPiece tokenizer (vocabulary size ~30k)	Byte-Pair Encoding (BPE) tokenizer (vocabulary size ~50k)

Training argument

```
training_args = TrainingArguments(  
    output_dir="/bert-review-classifier",  
    per_device_train_batch_size=16,  
    per_device_eval_batch_size=16,  
    num_train_epochs=4,  
    learning_rate=4e-5,  
    weight_decay=0.01,  
    logging_dir="/logs",  
    logging_steps=50,  
    save_total_limit=1,  
    save_strategy="epoch",  
    report_to="none",  
    fp16=True  
)
```

1.BERT-BASE

1.1 Dataset with stop word removal and lemmatization

Learning rate of 5e-5: Train accuracy=96%, val accuracy=58%

Learning rate of 4e-5: Train accuracy=95%, val accuracy=58%

Learning rate of 3e-5: Train accuracy=94%, val accuracy=59%

1.2 Dataset without stop word removal and lemmatization

Learning rate of 5e-5: Train accuracy=97% , val accuracy=63.4%

Learning rate of 4e-5: Train accuracy=97.1% , val accuracy=63.9%

Learning rate of 3e-5: Train accuracy=96.2% , val accuracy=63.7%

Learning rate of 2e-5: Train accuracy= 93.7%, val accuracy=64.3%

Learning rate of 1e-5: Train accuracy=84.2% , val accuracy=65.4%

2.ROBERTA-BASE

2.1 Dataset without stop word removal and lemmatization

Learning rate of 1e-5: Train accuracy=77% , val accuracy=67.4%

Learning rate of 2e-5: Train accuracy=87.1% , val accuracy=66.4%

Learning rate of 3e-5: Train accuracy=87% , val accuracy=66.4%

Result

Learning rate of 4e-5 without stop word removal and lemmatization using Bert base achieved higher accuracy of Train accuracy=97.1% , val accuracy=63.9%

Model was trained on whole dataset

```
from sklearn.metrics import classification_report
target_names = ["Rating 1", "Rating 2", "Rating 3", "Rating 4", "Rating 5"]
predictions_output = trainer.predict(new_dataset)
y_pred = np.argmax(predictions_output.predictions, axis=1)
y_true = predictions_output.label_ids

print("\nClassification Report:")
print(classification_report(y_true, y_pred, digits=4, target_names=target_names))
```

```
Classification Report:
              precision    recall  f1-score   support

 Rating 1       0.9270      0.9386      0.9328       27888
 Rating 2       0.8377      0.8137      0.8255       14823
 Rating 3       0.8673      0.8622      0.8648       21670
 Rating 4       0.8963      0.8897      0.8930       38052
 Rating 5       0.9380      0.9476      0.9428       47118

 accuracy              0.9056       149551
 macro avg       0.8933      0.8904      0.8918       149551
 weighted avg    0.9052      0.9056      0.9053       149551
```

The classification report summarizes the performance of the model across all five rating categories. For each rating (1 to 5), the metrics reported include **precision, recall, and F1-score**, along with the total number of instances (support).

- **Rating 1** achieved the highest balance between precision (0.9270), recall (0.9386), and F1-score (0.9328), indicating strong performance in correctly classifying this class.
- **Rating 2** shows the lowest performance overall, with a precision of 0.8377, recall of 0.8137, and F1-score of 0.8255, suggesting the model struggled more with this class compared to others.
- **Ratings 3, 4, and 5** show consistently high performance, with F1-scores ranging from 0.8648 to 0.9428, demonstrating that the model was effective at predicting these categories.

At the overall level, the model achieved an **accuracy of 90.56%**, with a **macro-average F1-score of 0.8918**, and a **weighted-average F1-score of 0.9053**. The macro-average indicates that performance was fairly balanced across classes, while the weighted-average accounts for class imbalance and confirms strong overall predictive ability.

Code:

```
model = AutoModelForSequenceClassification.from_pretrained('./epoch 6')
```

```
def compute_metrics(eval_pred):
```

```
    logits, labels = eval_pred
```

```
    predictions = np.argmax(logits, axis=-1)
```

```
    return {
```

```
        "accuracy": accuracy_score(labels, predictions),
```

```
        "f1": f1_score(labels, predictions, average="weighted")
```

```
    }
```

```
from transformers import TrainingArguments
```

```
training_args = TrainingArguments(
```

```
    output_dir='./results',
```

```
    per_device_eval_batch_size=16,
```

```
    do_train=False,
```

```
    do_eval=True,
```

```
    report_to="none",
```

```
    fp16=True
```

```
)
```

```
trainer = Trainer(
```

```
    model=model,
```

```
    args=training_args,
```

```
    tokenizer=tokenizer,
```

```
    compute_metrics=compute_metrics
```

```
)
```

```
from sklearn.metrics import classification_report
```

```
target_names = ["Rating 1", "Rating 2", "Rating 3", "Rating 4", "Rating 5"]
predictions_output = trainer.predict(new_dataset)
y_pred = np.argmax(predictions_output.predictions, axis=1)
y_true = predictions_output.label_ids

print("\nClassification Report:")
print(classification_report(y_true, y_pred, digits=4, target_names=target_names))
```